

# Deep Learning Lab6

110550085 房天越

## 1. Introduction

In this homework, we aim to implement a Denoising Diffusion Probabilistic Model to do Image Generation through multi-label condition. Then, use a pretrained model with ResNet18 on this task to test its accuracy, our goal is to make this score as high as possible.

## 2. Implementation Details

### 2.1 Model Architecture

We apply UNet2DModel from HuggingFace as the noise predictor.

The block channels are set as [128, 128, 256, 256, 512, 512].

The Down and Up blocks is a mix of standard and attention-based modules.

For the conditional embedding, we use one-hot labels to linear projection into a 512-dim embedding via `nn.Linear(num_classes, time_embed_dim)`, which we feed into UNet as class embeddings.

```
# UNet2DModel
self.unet = UNet2DModel(
    sample_size=sample_size,
    in_channels=3,
    out_channels=3,
    layers_per_block=2,
    block_out_channels=[128, 128, 256, 256, 512, 512],
    down_block_types=[
        "DownBlock2D",
        "DownBlock2D",
        "DownBlock2D",
        "DownBlock2D",
        "AttnDownBlock2D",
        "DownBlock2D",
    ],
    up_block_types=[
        "UpBlock2D",
        "AttnUpBlock2D",
        "UpBlock2D",
        "UpBlock2D",
        "UpBlock2D",
        "UpBlock2D",
    ],
    class_embed_type="identity",
)
```

### 2.2 Noise Scheduler

We employ DDPM Scheduler with 1000 training timesteps and the `squaredcos_cap_v2` schedule, which is an improved cosine schedule for smoother noise variance decay.

```
def __init__(
    self,
    num_classes: int = 24,
    time_embed_dim: int = 512,
    sample_size: int = 64,
    beta_schedule: str = "squaredcos_cap_v2",
    guidance_scale: float = 2.0,
):
    super().__init__()
    self.guidance_scale = guidance_scale

    # Scheduler for noise
    self.scheduler = DDPMSScheduler(
        num_train_timesteps=1000,
        beta_schedule=beta_schedule,
        prediction_type="epsilon",
    )
```

### 2.3 Classifier-Free Guidance

During training, we drop the class embedding to simulate unconditional generation with probability = 0.2.

At sampling time, we compute:

$$\epsilon = \epsilon_{\text{uncond}} + s \cdot (\epsilon_{\text{cond}} - \epsilon_{\text{uncond}})$$

, where s is the guidance scale we can adjust.

```
def sample(self, y, device, num_inference_steps=1000):
    b = y.shape[0]
    x = torch.randn(b, 3, 64, 64, device=device)
    scheduler = self.scheduler
    scheduler.set_timesteps(num_inference_steps, device=device)

    y_zero = torch.zeros_like(y)
    for t in scheduler.timesteps:
        eps_uncond = self.unet(x, t, self.class_embedding(y_zero)).sample
        eps_cond = self.unet(x, t, self.class_embedding(y)).sample
        eps = eps_uncond + self.guidance_scale * (eps_cond - eps_uncond)
        step = scheduler.step(eps, t, x)
        x = step.prev_sample

    return x

def add_noise(self, x, noise, timesteps):
    return self.scheduler.add_noise(x, noise, timesteps)
```

### 2.4 Training Configuration

Data Loader: A custom IClevrDataset reads JSON for image mappings, resize to 64\*64, converts to tensor, and normalizes. The batch size is set to 32.

```

class IClevrDataset(Dataset):
    def __init__(self, json_path, image_dir, transform=None):
        with open(json_path) as f:
            self.conditions = json.load(f)
            self.image_dir = image_dir
            self.transform = transform or transforms.Compose([
                transforms.Resize((64,64)),
                #transforms.RandomHorizontalFlip(p=0.5),
                #transforms.ColorJitter(0.1,0.1,0.1,0.1),
                transforms.ToTensor(),
                transforms.Normalize((0.5,)*3, (0.5,)*3)
            ])
            self.keys = list(self.conditions.keys())

    def __len__(self):
        return len(self.keys)

    def __getitem__(self, idx):
        fname = self.keys[idx]
        labels = self.conditions[fname]
        path = os.path.join(self.image_dir, fname)
        img = Image.open(path).convert('RGB')
        x = self.transform(img)
        y = torch.zeros(len(OBJECTS_MAP))
        for obj in labels:
            y[OBJECTS_MAP[obj]] = 1
        return x, y

```

Optimizer and Scheduler: Adam with  $1e-5$  and CosineAnnealingLR over the epochs.

Mixed Precision: We use torch.cuda.amp to accelerate training and save memory.

```

optimizer = optim.Adam(model.parameters(), lr=args.lr)
lr_scheduler = optim.lr_scheduler.CosineAnnealingLR(optimizer, T_max=args.epochs)
scaler = GradScaler() # AMP

```

Checkpoint and Logging: We save model weights after each epoch and log loss/learning-rate scalars with TensorBoard.

```


lr_scheduler.step()
writer.add_scalar('train/loss', np.mean(losses), epoch)
writer.add_scalar('train/lr', lr_scheduler.get_last_lr()[0], epoch)

ckpt_path = os.path.join(args.out_dir, f'ddpm_epoch{epoch}.pth')
torch.save({'model': model.state_dict()}, ckpt_path)


```

### 3. Results and Discussion


#### 3.1 test.json



### 3.2 new\_test.json



### 3.3 Denoising Image of ["red sphere", "cyan cylinder", "cyan cube"]



### 3.4 Discussion

#### 3.4.1 Guidance Scale

During the training process, we compute:

$$\epsilon = \epsilon_{\text{uncond}} + s \cdot (\epsilon_{\text{cond}} - \epsilon_{\text{uncond}})$$

I tried to use  $s = 1$  (Fully conditional),  $s = 1.5$ , and  $s = 2$ , and run for 100 epochs. No significant difference in accuracy is observed.

#### 3.4.2 Noise Scheduler

I tried to use linear, cosine, and squaredcos\_cap\_v2 schedulers, the performance is squaredcos\_cap\_v2 > cosine > linear.

The accuracy of linear is about 0.6 after 200 epochs.

The accuracy of cosine is about 0.7 after 200 epochs.

The accuracy of squaredcos\_cap\_v2 is around 0.88 (tested many times) after 200 epochs.

## 4. Experimental Results

```
test set (./test.json) accuracy: 84.72%
```

```
new_test set (./new_test.json) accuracy: 90.48%
```